

Unit 13
Big Data Analytics

Characteristics of Many Modern OLAP Applications

- ◆ Datasets keep increasing in size
- ◆ Individual nodes can not process entire dataset efficiently
- ◆ Individual nodes can not store entire dataset efficiently
- ◆ Need to resort to distributed processing and storage
- ◆ For storage: distributed filesystems
 - E.g., HDFS with super-fast reads but only append writes
- ◆ Luckily for many applications “reading” is sufficient
 - No writing to datasets as part of application
 - Performed in a second step, separate system, if needed

Approaches

- ◆ Many approaches exist for storing and/or processing data
- ◆ Key is **efficient mapping** of data to underlying hardware infrastructure based on 1. structure of data, 2. functionalities offered, and 3. guarantees provided, e.g.,
 - 1: OO vs relational data vs (*key*, *value*) pairs
 - 2: arbitrary position read&write vs arbitrary position read&append only vs sequential read&append only
 - 3: isolation for sequence of operations vs individual operations only
- ◆ Gives rise to **many** different “big data” approaches and systems, e.g., MapReduce, Pig/PigLatin, Flume, Mahout, Hive, Samza, Apex, Ignite, Kafka, Hedwig, Storm, Heron, Flink, Spark
 - **Understanding tradeoffs (1./2./3.) and system concepts is crucial for making the right choice**
 - Vs pure SE principles (e.g. lines of code, modularization)
- ◆ We will briefly discuss three:
 1. MapReduce
 2. Spark/RDDs
 3. Storm

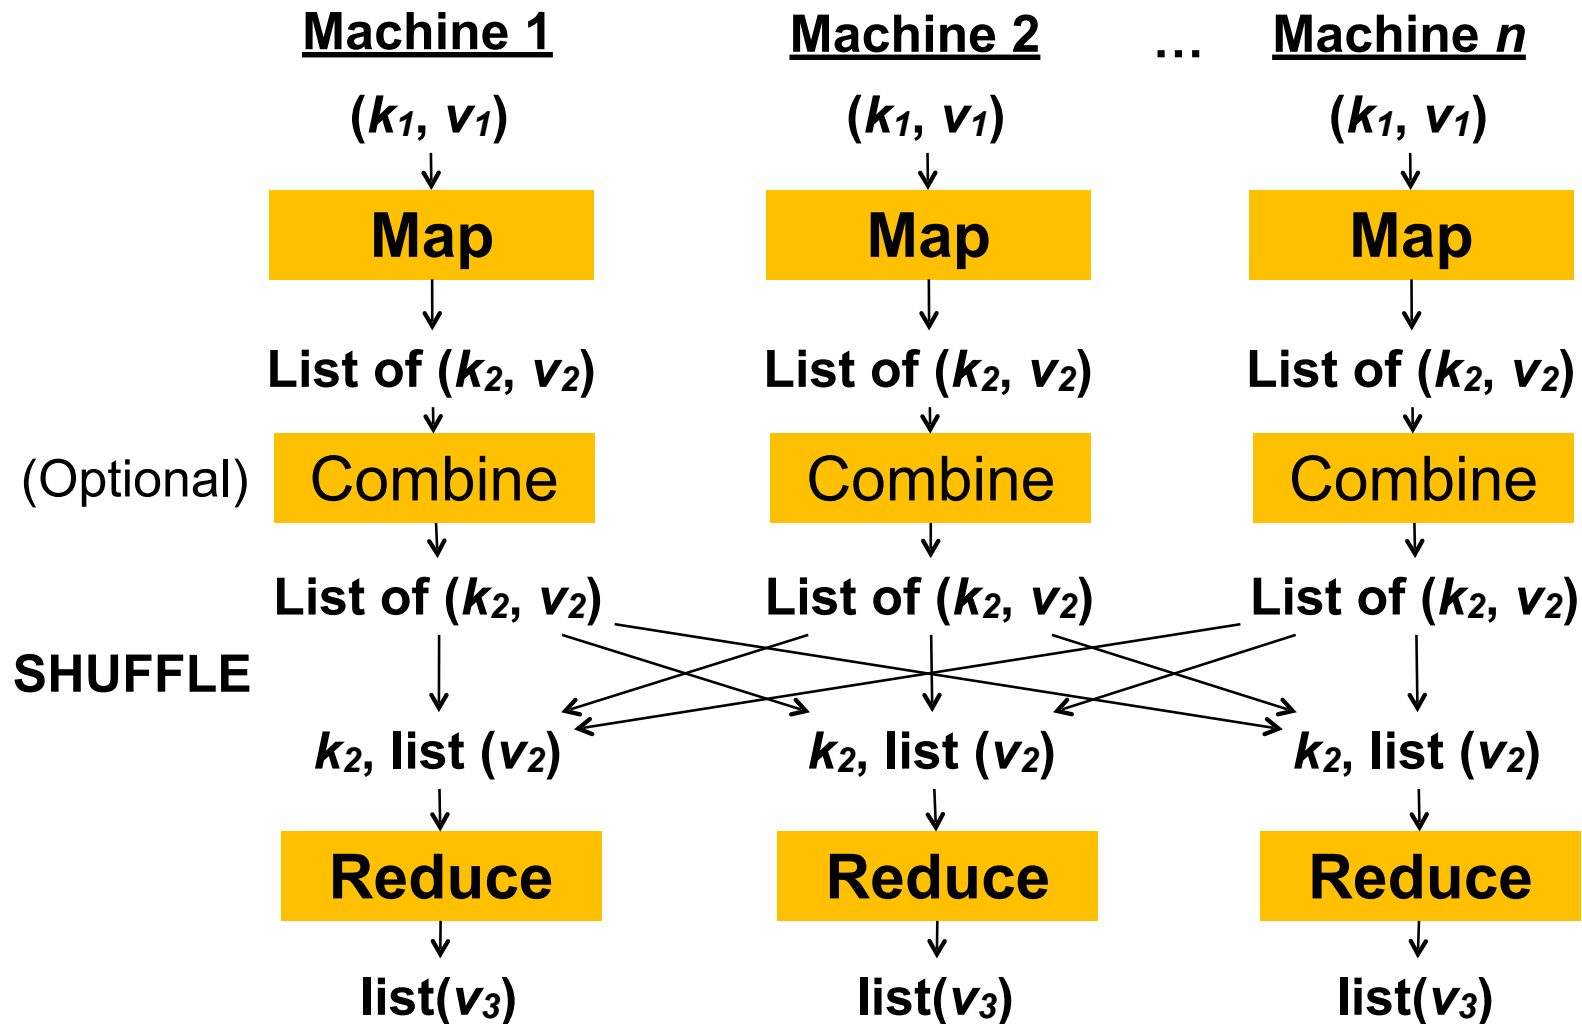
MapReduce (a.k.a. Map/Reduce, Map-Reduce, ...)

- ◆ Introduced originally in Lisp programming language
- ◆ Popularized for distributed processing by Google in 2004 [Dean&Ghemawat;OSDI'04]
- ◆ Allows for distribution with strong potential for parallelization
- ◆ Well-suited for data structured as (*key*, *value*) pairs
- ◆ Simple abstraction, programmer implements 2 functions:
 - $map(k_1, v_1) \rightarrow list(k_2, v_2)$
 - $reduce(k_2, list(v_2)) \rightarrow list(v_3)$
- ◆ Note: k_2 s needn't be k_1 s, v_2 s needn't be v_1 s, v_3 s needn't be v_2 s

Workflow

1. Data loaded in from file, n partitions created
 2. n parallel mappers instantiated on n hosts, each obtains a partition
 3. *map* function called for one partition entry (k_1, v_1) pair at a time, returns list of (k_2, v_2) pairs each time
 4. Temporary files created with values v_2 per (same) key k_2
 5. m parallel reducers instantiated on m hosts
 6. Each reducer is assigned a subset of the keys k_2
 7. *reduce* function called for a key k_2 at a time with *all* values from temporary files created for that key *by any* mapper, outputs list of values v_3 each time
- ◆ System phase taking care of 6. and distributing data to reducers is called “shuffling”
 - ◆ System also responsible for monitoring “stragglers” (slow mappers or reducers) and re-initiating them in case of failures, as well as load balancing across hosts

Schematic Overview (Special Case of $m=n$)



Example: Word Count (Pseudo-Code)

- ◆ Input: set of files/documents
- ◆ Output: words and the number of their respective occurrences

```
map(String input_key, String input_value):
```

```
    // input_key: document name
```

```
    // input_value: document contents
```

```
    for each word w in input_value:
```

```
        EmitIntermediate(w, "1");
```

```
reduce(String output_key, Iterator intermediate_values):
```

```
    // output_key: a word
```

```
    // output_values: a list of counts
```

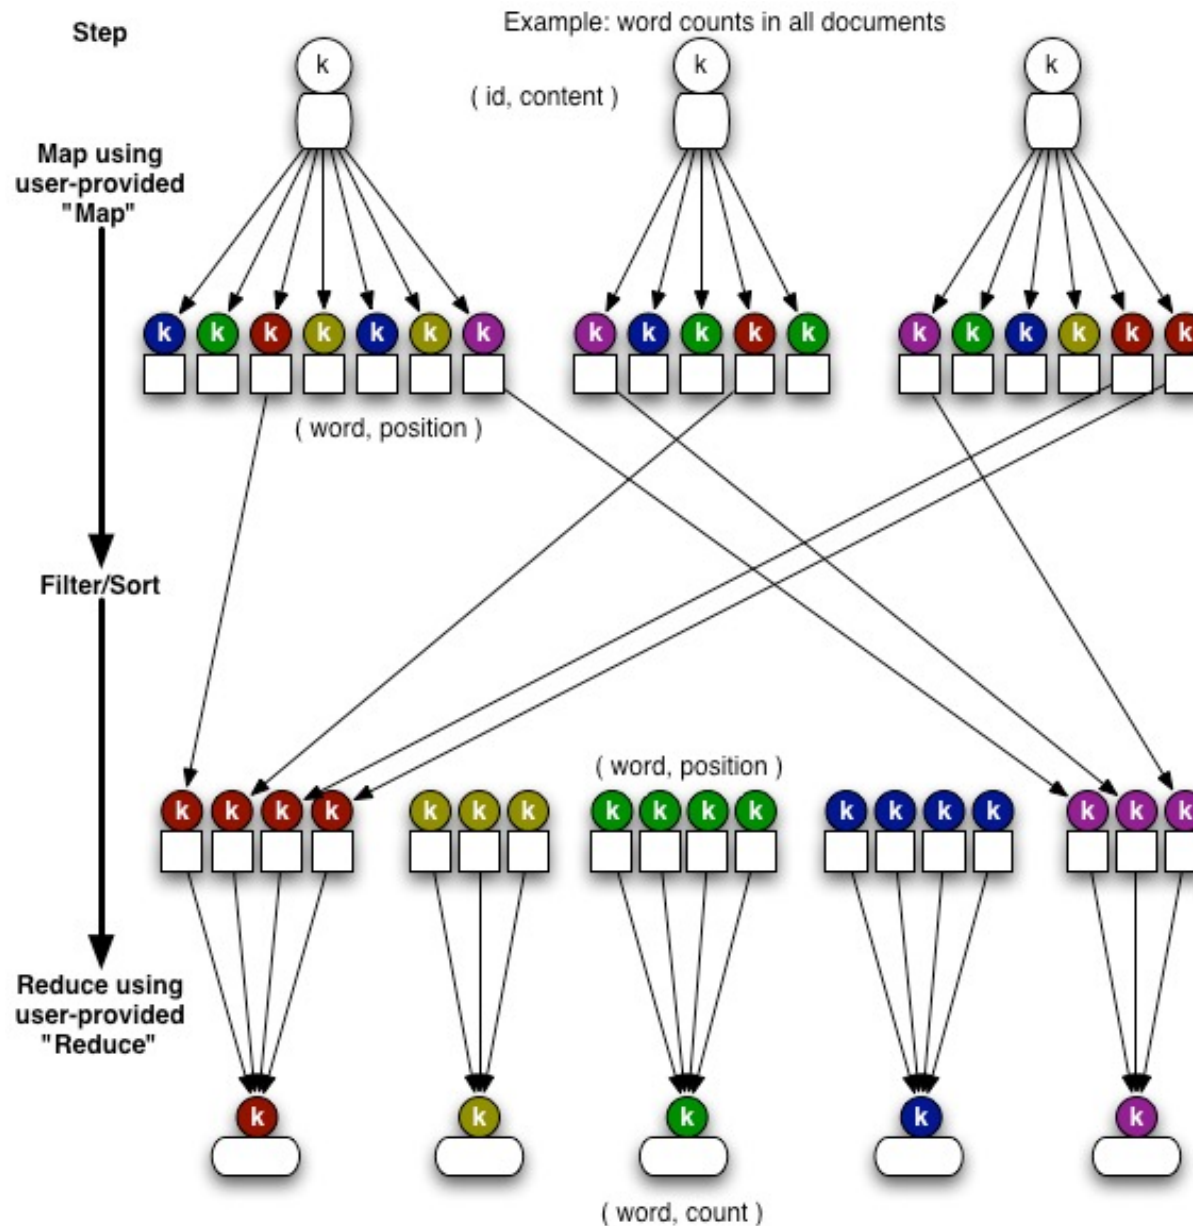
```
    int result = 0;
```

```
    for each v in intermediate_values:
```

```
        result += ParseInt(v);
```

```
    Emit(AsString(output_key + ":" + result));
```

Word Count Illustrated



Limitations

◆ Model

- Not all (parallelizable) tasks fit the abstraction
- Can chain multiple **stages** with respective *map* and *reduce* functions
 - Still no unbounded iterative (and recursive) computation stages (cf. great^x-grandchildren/parents)

◆ Performance

- Filesystem becomes bottleneck
- Can avoid distributed filesystem between different stages, but shuffling still uses local filesystem

◆ Many extensions proposed (for model and/or performance), e.g., aggregators (reducer-side)

Spark

- ◆ Original paper published in 2012 [Zaharia et al.;NSDI'12]
- ◆ Key tenets
 - A. **Dataflow-based** computation from input data sets to results through sequence of operations
 - Main abstraction: **resilient distributed datasets (RDDs)**
 - No modifications in place, modifications give rise to new RDD
 - Computation forms a DAG with nodes RDDs and arcs for operations
 - B. **Avoid filesystem** during computations
 - All datasets **materialized** in RAM including **intermediate** datasets
 - Support for **iterative** and **incremental** computations
 - C. Scaling by **partitioning datasets across hosts** (RAM)
 - API highlights operations that can be performed on individual partitions
 - D. **Lineages** for tracking dependencies between datasets
 - Used for incremental computations and fault-tolerance (recomputing)
- ◆ A. from data flow programming (e.g., PigLatin [Olsen et al.;SIGMOD'08]), B. from main-memory DBMSs (e.g., [Pedone&Frolund;SRDS'00]), C. from distributed DBMSs, D. from snapshot techniques and provenance tracking (e.g., [Zhang et al.;VLDB'07])

Operations on RDDs

1. Transformations

- Create RDDs from other (parent) RDDs
- Typically one element at a time (and thus independently across partitions)
- Some common transformations
 - *map(function)*: *function* applied to every element to create new element
 - *filter(function)*: *function* applied to every element to decide whether to “keep” it

2. Actions

- Compute actual return values
- Some common actions
 - *count()*: return the number of elements
 - *take(n)*: return an array of first *n* elements
 - *collect()*: return array of all elements
 - *saveAsTextFile(file)*: save to *file* in filesystem
 - *reduce(function)*: aggregate all values using *function*

◆ Many alternative APIs and libraries, e.g.,

- DataFrame (tables w/ named columns), Dataset (static typing for OOP)
- SparkSQL
- Spark streaming (discretized streams)
- SparkML/MLlib (on DataFrame)

Examples (Pseudo-Code)

◆ Word count

```
counts = sc.textfile("hdfs://...")  
        .flatMap(lambda line: line.split('\s'))  
        .map(lambda word: (word, 1))  
        .reduceByKey(operator.add)  
counts.save("hdfs://...")
```


◆ General MapReduce

```
result = data.flatMap(map_fn)  
           .groupByKey()  
           .map(lambda (k, vs): reduce_fn(k, vs))
```

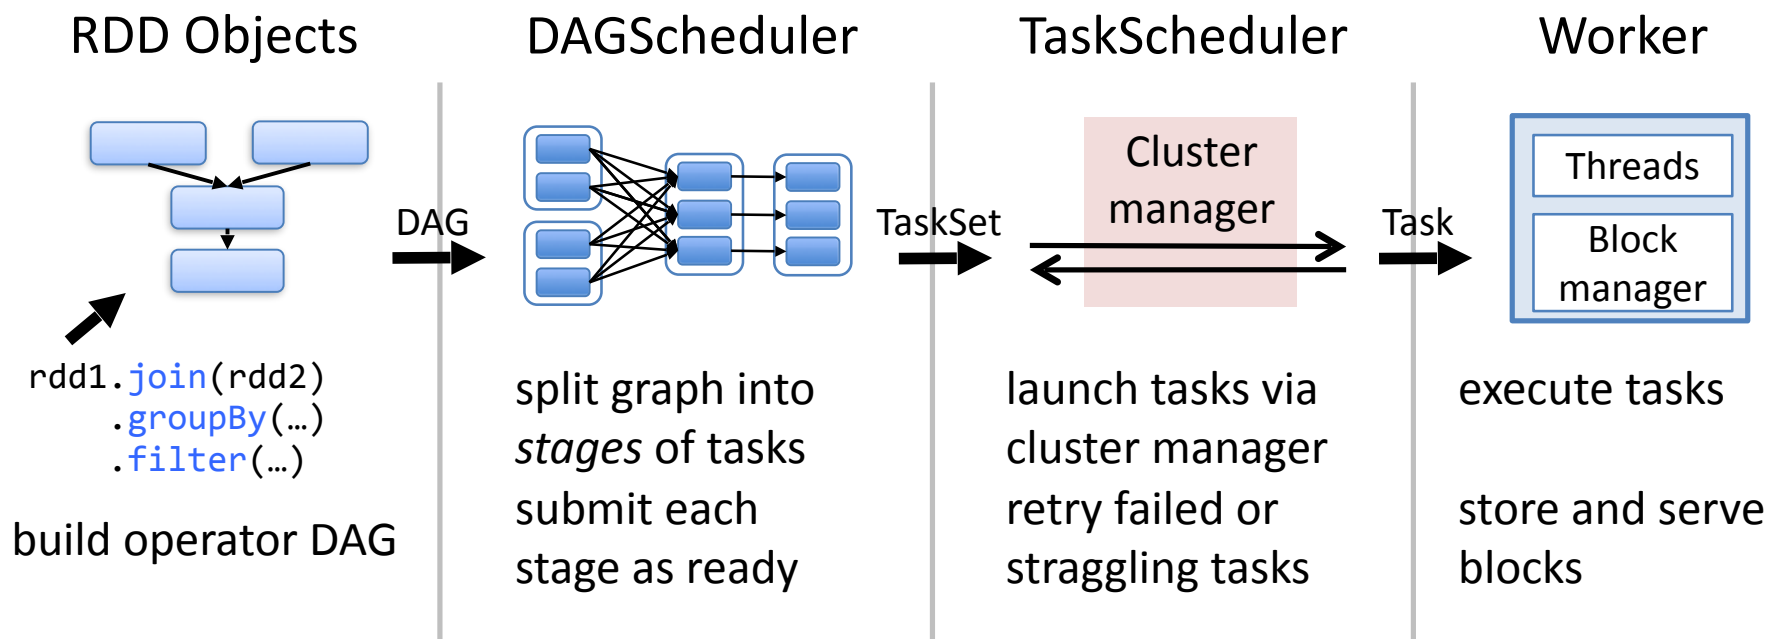
◆ General MapReduce with combiner

```
result = data.flatMap(map_fn)  
           .reduceByKey(combiner_fn)  
           .map(lambda (k, vs): reduce_fn(k, vs))
```


Architecture Overview



Tracing Execution



MapReduce vs Spark: Logistic Regression



Limitations

◆ Aggregation

- Performance excellent when “embarrassingly parallel”
- Aggregation performance depends on aggregation function used
 - Sometimes best in many stages $2 \times 2 \times 2 \dots$ (e.g., HP Presto [Venkataraman et al.;EuroSys’13])
 - Sometimes best in 1 stage $n \times 1$ (e.g., Spark original)
 - Often in between, depending on **aggregation ratio** (tradeoff computation vs communication)
- Cf. LOOM [Culhane et al.;HotCloud’14]*,[Culhane et al.;INFOCOM’15], ROME [Blöcher et al.;ACM TOCS’22]
 - `treeReduce` (added after *) only poor subset solution using **depth** parameter

◆ Security

- Cf. Seabed [Papadimitriou et al.;OSDI’16], Cuttlefish [Savvides et al.;SoCC’17], Opaque [Zheng et al.;NSDI’17], Symmetria [Savvides et al.;VLDB’20], Hydra [Mangipudi et al.;PLDI’23]

◆ Correctness

- Cf. Multi-party session types for event-driven fault-tolerant distributed programming [Viering et al.;OOPSLA’21]


...

Another Issue: Reactivity

- ◆ Spark et al. are optimized for batch/mini-batch processing
- ◆ Support for incremental computation upon extension of input datasets
- ◆ What if arrival of new data is the norm and reaction times need to be minimized?
- ◆ Enter ***stream processing***
- ◆ Scenarios
 - Financial applications (e.g., electronic/algorithm trading, fraud detection)
 - Network monitoring
 - Social network analysis
 - Sentiment analysis on tweets
 - ...
- ◆ Stream processor listens indefinitely for incoming data
 - Reacts

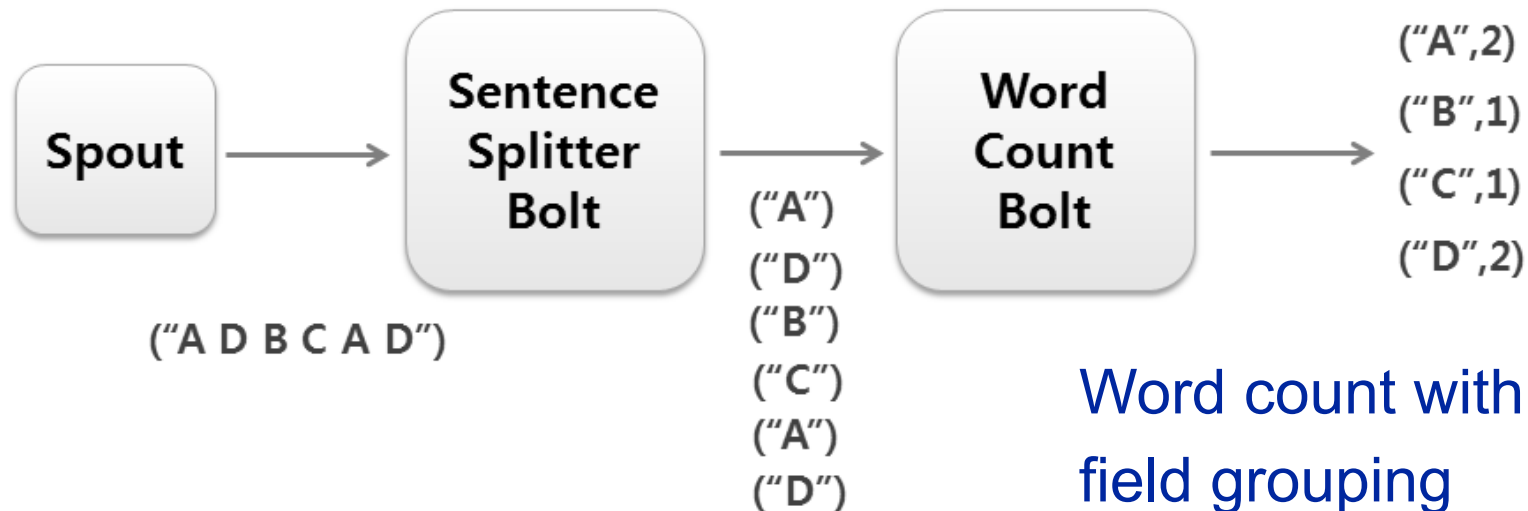
Storm

- ◆ Relies on explicit application-defined topology
- ◆ Concepts
 - **Tuple**: ordered list of named values
 - **Stream**: an unbounded sequence of tuples
 - **Spout**: source of a stream emitting tuples
 - **Bolt**: accepts tuple from (one if its) input streams, performs some computation (filtering, aggregation, join), possibly emits new tuple(s)



Parallelism

- ◆ Bolts can be replicated
- ◆ **Groupings** specify how tuples are routed to bolt replicas
 - Shuffle grouping: random distribution
 - Field grouping: portioning according to value of tuple attribute
 - All grouping: complete replication
 - Direct grouping: producer decides on replica
 - ...
 - Custom application-defined grouping



Key Ideas

- ◆ Distributed data analytics
- ◆ MapReduce
- ◆ Spark
- ◆ Limitations/open issues
- ◆ Stream processing